Lecture 3 — Relational Query Languages and SQL

Jeff Zarnett

jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

September 12, 2023

A query language is a language to get information from the database.

"Excuse me, I'd just like to ask a question..."

Procedural: you tell the database exactly how to get the data you want.

Non-procedural: tell the database what you want; it figures out how to deliver.

Operations can be applied to a single relation, or a pair of relations.

The result of this is always a single relation.

This means that the output of one operation can be used as the input to another operation, allowing us to chain operations as needed.

If it helps you to imagine this, every operation is a function with return type `Relation` and they take either one or two `Relations` as parameters.

You will also need to keep in mind that because they all have input and output of relation(s), the operations operate on a set of data, not just a single element.

In a C-like language, you write a `for` loop to sets all elements in an array to -1.

In SQL, you write a query that sets the value to -1 for all entries in the relation.

This means that if your normal mode of thinking is C-like (procedural) then there is a mental transition that needs to be made.

FOR THAT ONE FRACTION OF A SECOND, YOU WERE OPEN TO OPTIONS YOU HAD NEVER CONSIDERED. THAT IS THE EXPLORATION THAT AWAITS YOU.

The fundamental operations in relational algebra are:

- Select
- Project
- Union
- Set Difference
- Cartesian Product
- Rename

And there are three shortcuts we will use that can be created using the fundamental operations:

- Set Intersection
- Assignment
- Join

SQL – "Structured Query Language" – is not only for querying but also for defining and changing the database schema.

It is a nonprocedural language: you ask for what you want and the database server returns a result that matches that format.

Other query languages do exist.

Reality check: the majority of students taking this course already have exposure to some form of database system.

Most likely a SQL relational database of some sort.

That means that a course where we spend a lot of time on SQL is redundant. But we can't skip it either... It can still have value.

Still, the mathematical notation may be new.

- Data Definition Language (DDL)
- Data Manipulation Language (DML)
- Data Control Language (DCL)
- Transaction Control Language (TCL)

VIN	year	make	model	license_plate_number
2B4FH25K1RR646348	2005	Honda	Civic	ZZZZ 249
4UZACLBW87CZ42980	2011	Ford	Focus	YYYY 995
WVWDM7AJ4BW227648	2015	Volkswagen	Golf	ZRHC 112
1J4FT38L7KL506678	2017	Audi	A4	VNTY PLT
2C8GF48415R447850	2016	Volkswagen	Jetta	VWJG 071
2T3RK4DV1AW089914	2014	Toyota	Camry	ZZZZ 385
1XPCDR9X2YN436206	2012	Toyota	Camry	ZZYY 251
1GYS3BKJ5FR338462	2016	Honda	Civic	YYYY 585
1GAGG25R6Y1243081	2015	Chevrolet	Corvette	GORLFAST
JH2SC59208M054959	2018	Genesis	Genesis	AAAA 123

The first operation we will discuss is selection: find the tuples in the relation.

And bring me the passengers. I want them alive!

We need to specify what we want to find, and where we want to look for it.

We can specify that we want to find only the things that match certain criteria.

In mathematical notation, selection has the symbol σ (sigma).

The predicate is put in a subscript and then the function parameter, the relation to select from, is follows in parenthesis.

$\sigma_{make=''Volkswagen''}(vehicle)$:

VIN	year	make	model	license_plate_number
WVWDM7AJ4BW227648	2015	Volkswagen	Golf	ZRHC 112
2C8GF48415R447850	2016	Volkswagen	Jetta	VWJG 071

We specify the predicate *p* using propositional logic.

A term is written as an attribute followed by a comparison operator then an attribute or constant.

The comparison operators are $=, \neq, >, \geq, <, \leq$.

So the attribute in the example is *make* and the comparison operator is $=$ (equals) and then the right hand side is the constant "Volkswagen".

Terms are connected by the mathematical AND $\wedge$, mathematical OR $\vee$, and mathematical NOT $\neg$ operators.

Thus we could specify make = "Volkswagen" and year = "2015":

$(\sigma_{make=''Volkswagen'' \wedge year=''2015''}(vehicle))$:

VIN	year	make	model	license_plate_number
WVWDM7AJ4BW227648	2015	Volkswagen	Golf	ZRHC 112

Compound Predicate

We can chain as many of these qualifiers as desired.

Use of parenthesis is recommended to make it clear how the boolean logic is to be evaluated.

If you wish the year to be either 2015 or 2016 and the make to be Volkswagen:

$make =''$ *Volkswagen''* $\land$ $(year =''$ $2015''$ $\lor$ *year* $=''$ $2016'')$.

It doesn't cost extra money to write the parenthesis so use them anywhere there is the possibility of error or confusion.

It is possible that no tuples in the relation match all of the predicates.

If the predicate is make = "Chrysler" then no matches will be found.

If we specify make = "Honda", and model = "Camry", we will still find no tuples that match that either. Honda doesn't manufacture any model named "Camry".

Either way we get back an empty relation.

In SQL, the `SELECT` clause is used to, well, select tuples from the database.

The SQL command corresponding to $\sigma(vehicle)$ is `SELECT * FROM VEHICLE;`

It is traditional in SQL that we write keywords in all capitals.

A statement is terminated by a semicolon.

Note the use of FROM to specify the relation.

The asterisk (∗) is needed to specify what we are going to select.

The select statement contains no predicates so it would return all tuples in the relation.

```
SELECT * FROM VEHICLE WHERE make = 'Volkswagen';
```

The WHERE keyword indicates the start of the predicate.

The string literal "Volkswagen" is enclosed in single quotation marks but it can also be in double quotation marks

To form a compound predicate: SELECT * FROM VEHICLE WHERE make = 'Volkswagen' AND year = '2015';.

SQL Predicate Comparison

The comparison operators <, >, >=, <= all look like C.

Equals is just a single equals sign (=).

The not equals operator is written <>.

These operations can work on strings, mathematical types, dates, et cetera.

You may get unexpected results if you attempt to compare two types that don't compare well.

It is possible to chain the where clause to include a range, such as `WHERE year >= 2010 AND year <= 2015`.

there is a bit of notational convenience in SQL: the `BETWEEN ... AND` syntax.

Thus one can write `WHERE year BETWEEN 2010 AND 2015`.

This is inclusive on both ends, so be careful about whether this is the desired behaviour.

The predicate may also contain the name of other attributes in the relation.

```
SELECT * FROM VEHICLE WHERE make = model;.
```

This returns the 1 tuple in the sample data where make and model are the same.

There are situations in real life where we want to use an attribute rather than a constant, although in this particular example it may look a little strange.

The second operation is projection.

This takes a subset of a relation: take a relation and extract from it only the things that we asked for.

Projection cuts down the relation to the specific attributes.

Projection takes a single relation and returns another relation.

If we want to know what the VIN is for the car with the license plate "GORLFAST" we can ask for just that and need not get back the entire tuple.

Similarly, if you want to get a list of the makes and models in the database, again, it's nice to be able to get this data without any parts you do not need.

Projection also makes certain things a lot faster: it makes no sense to load all this extra data to your program just to ignore almost all the columns.

In mathematical notation, projection has the symbol Π (capital pi).

Like selection, it takes a subscript and the input relation follows in parenthesis.

So, $\Pi_{make,model}(vehicle)$ will reduce the relation down to just the attributes make and model, leaving out VIN, year, and license plate number information.

All tuples in the relation will appear in the output of the projection operation, although it is entirely possible that some of them look very much alike.

make	model
Honda	Civic
Ford	Focus
Volkswagen	Golf
Audi	A4
Volkswagen	Jetta
Toyota	Camry
Chevrolet	Corvette
Genesis	Genesis

There are fewer tuples in this relation than there were in the original relation.

Remember: a relation is a set and therefore duplicates are not permitted.

The order of the tuples is not specified either.

The attributes will be listed in the order they are provided in the project clause, which we may choose arbitrarily.

In SQL projection is done through our application of the SELECT clause.

In the previous examples we used ∗ which specified all attributes.

In production code, the use of SELECT ∗ is discouraged.

`SELECT make, model FROM VEHICLE;` returns:

make	model
Honda	Civic
Ford	Focus
Volkswagen	Golf
Audi	A4
Volkswagen	Jetta
Toyota	Camry
Toyota	Camry
Honda	Civic
Chevrolet	Corvette
Genesis	Genesis

Detection of duplicates is relatively difficult (computationally expensive).

SQL does not remove them unless you ask, so technically what you get back is more akin to a list than to a set.

If you wish to remove duplicates you may use the keyword `DISTINCT`:
`SELECT DISTINCT make, model FROM VEHICLE;`.

The opposite of `DISTINCT` is the `ALL` keyword, but since this is the default behaviour you are unlikely to need to write it.

The SQL `SELECT` statement is then both selection and projection.

In what order do these operations take place? Sometimes it does not matter, e.g., if we asked for `SELECT *...` then the projection does nothing.

Here's a hint: this statement is perfectly valid: `SELECT make, model from VEHICLE WHERE year = '2016';` and returns:

make	model
Volkswagen	Jetta
Honda	Civic

Clearly the selection operation must take place first!

If we cut the attributes down to just make and model, then we lose the information we need to tell whether the year is 2016 or not.

The `SELECT` statement in SQL lets us do some interesting things other than select attributes alone.

The select clause may contain arithmetic expressions using the operators $+, -, *, /$ operating on either attributes or tuples.

Thus, to get a 2 digit date out of a 4 digit date we could do: `SELECT year - 2000 FROM VEHICLE;`.

We can also select constants: write that constant instead of the name of the attribute.

```
SELECT 'Car', make, model FROM vehicle WHERE year = '2016';
```
produces:

Car	make	model
Car	Volkswagen	Jetta
Car	Honda	Civic

We can combine two relations using the union operation.

It takes two relations and produces a single relation from them.

The mathematical symbol is ∪, the same as the mathematical set union operator.

Typically this is used in conjunction with other operations (selection, projection) to get the data that you want.

The query $\sigma_{make=''Ford''}(vehicle) \cup \sigma_{make=''Audi''}(vehicle)$ produces the results:

VIN	year	make	model	license_plate_number
4UZACLBW87CZ42980	2011	Ford	Focus	YYYY 995
1J4FT38L7KL506678	2017	Audi	A4	VNTY PLT

In SQL the keyword is `UNION`.

Use of parenthesis is recommended (if not mandatory) to help the SQL query parser figure out what it is you are asking for.

```
(SELECT license_plate_number FROM vehicle WHERE make = 'Ford')
UNION (SELECT license_plate_number FROM vehicle WHERE make =
'Audi');
```

By why not write this as a compound predicate?

Union is for when we want to combine data from two different relations.

$(\Pi_{name,street,city,province,postal\_code}(\sigma(owner\_address))) \cup$
$(\Pi_{name,street,city,province,postal\_code}(\sigma(employee)))$

In SQL:

```
(SELECT name, street, city, province, postal_code FROM
owner_address)
UNION
(SELECT name, street, city, province, postal_code FROM
employee);
```

Unite!

The union operation can only succeed if the types are compatible.

For the union operation to make sense, two conditions must hold:

1. The relations must have the same number of attributes.
2. The domain of attribute i in the first relation must be the same as the domain of attribute i in the second relation, for all i.

This example also is a situation where we might make use of a selection with a constant.

If we wanted to produce a relation with all addresses, and wish to include some type information, we could do this:

```
(SELECT 'Vehicle Owner', name, street, city, province,
postal_code FROM owner_address)
UNION
(SELECT 'MTO Employee', name, street, city, province,
postal_code FROM employee);
```

	name	street	city	province	postal_code
Vehicle Owner	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
Vehicle Owner	Thomas Anderson	1234 Main St	Waterloo	ON	A0A 0A0
Vehicle Owner	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
MTO Employee	Kim Morgan	491 Example Road	Waterloo	ON	N2N 2N2
MTO Employee	Jordan Singh	24 Eastern Avenue	London	ON	N6A 0B0