

Lecture — Relational Query Languages & SQL, Continued

Jeff Zarnett
jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

August 11, 2023

The **set difference** operation works more or less as you would expect: it takes a relation and removes from it any tuples that are in the second relation.

After taking away things we do not want, then we are left with the things we do want.

Its input is therefore two relations and it produces a new relation as its output.

The mathematical symbol is $-$ (the minus or subtraction operator).

The notation is then $r_1 - r_2$ where r_1 and r_2 are relations.

This produces a relation r_3 that contains all tuples in r_1 that are not in r_2 .

It is possible to chain the subtraction operators, but like regular mathematical subtraction, it does not commute: the order matters a lot.

The query $\sigma_{make="Honda"}(vehicle) - \sigma_{year < "2010"}(vehicle)$ produces the results:

VIN	year	make	model	license_plate_number
1GYS3BKJ5FR338462	2016	Honda	Civic	YYYY 585

Just like the union operation, the sets must be compatible. The same rules from union apply:

- 1 The relations must have the same number of attributes.
- 2 The domain of attribute i in the first relation must be the same as the domain of attribute i in the second relation, for all i .

In SQL the keyword we need for this is EXCEPT¹.

Like the union operation, parenthesis prevent confusion and/or make your database server less sad.

```
(SELECT make, model FROM vehicle WHERE make = 'Volkswagen')  
EXCEPT  
(SELECT make, model, FROM vehicle WHERE year < 2016);
```

Produces as output:

make	model
Volkswagen	Jetta

¹Although Oracle uses MINUS instead.

Those with experience in databases may also think about NOT IN.

The EXCEPT keyword encompasses both the DISTINCT and NOT IN behaviour, so duplicates are eliminated.

Example:

```
SELECT make, model FROM vehicle WHERE make = 'Volkswagen' AND  
VIN NOT IN  
(SELECT VIN FROM vehicle WHERE year < 2016);
```

This demonstrates a subquery: the first query selects a relation with one attribute.

The **cartesian product** combines information from two relations.

It is often the case that the data we need is in more than one relation.

That might be a design problem, but good design means data is in there only once and sometimes we do need to combine relations.

Suppose we wanted to send reminder notices to people whose license plates will expire next month. We combine license plate data with address data.

The mathematical symbol for cartesian product is \times , the multiplication symbol.

The cartesian product $r_1 \times r_2$ forms a third relation r_3 .

The new relation has all the attributes of both the relations that went into it (in the order specified).

Cartesian Product Example

Let us pretend our license plate relation contains exactly two tuples:

number	expiry	owner_address_id
ZZZZ 123	2018-09-30	24601
AAAA 855	2019-04-01	12949

And our owner address set is also two tuples:

id	name	street	city	province	postal_code
24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2

Cartesian Product Example

Then the query *license_plate* \times *owner_address* produces:

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
ZZZZ 123	2018-09-30	24601	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
AAAA 855	2019-04-01	12949	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2

This is all possibilities!

Cartesian Product's Cardinal Problem

Each tuple from the first relation is paired with each tuple from the second.

So if r_1 contains x tuples and r_2 contains y tuples, there will be $x * y$ tuples in the resulting relation...

The vast majority of which are garbage!

If we had two names that were the same, we would need a way to differentiate them.

Suppose that two relations, one called `book` and one called `author`, each having an attribute `id`.

Prefix the attribute with the name of the relation from which it came.

Thus, in the cartesian product the two columns would be `book.id` and `author.id`.

In SQL it is very simple to get the cartesian product:

```
SELECT * FROM license_plate, owner_address;
```

Separate the relations we wish to appear in the cartesian product with a comma.

We can always combine more if we needed.

We need to restrict our query with a selection predicate.

If the two relations are connected by some ID in in common, we need that.

In this case, the `id` attribute in the address matches `owner_address_id` in the license plate relation.

So $\sigma_{owner_address_id=id}(license_plate \times owner_address)$.

Or in SQL: `SELECT * FROM license_plate, owner_address where owner_address_id = id;`

This may seem unclear so we might prefer to prefix these with their names:
`SELECT * FROM license_plate, owner_address where license_plate.owner_address_id = owner_address.id;`

Adding the Restriction to Cartesian Product

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2

The **rename** operation is used to both change the name of existing attributes/relations and assign names to ones that have no name.

The mathematical notation is ρ (rho).

If applied to a relation, it is $\rho_x(E)$ where it renames the relation to x .

Or it can be applied to the name of the attributes: $\rho_{x(a_1, a_2, a_3, a_4 \dots)}(E)$ which renames the relation to x and then the attributes $a_1 \dots$

In SQL we can use the AS keyword.

That is not super interesting on a simple query to rename a relation:

`SELECT * FROM vehicle AS autos;` renames the relation result to “autos”.

It makes some more sense if we are going to use it in a subquery it will change the prefix for any duplicate attribute names (e.g., change `book.id` to `textbook.id`).

A more likely use is to rename the attributes of a relation.

The motivation for the rename operation right now probably seems weak.

There are some situations where we will need to use the rename operation.

Set intersection is exactly what it sounds like based on our understanding of mathematical sets.

The symbol for it is \cap and it is equivalent to $r_1 - (r_1 - r_2)$.

In SQL the keyword is INTERSECT.

Its use is limited, however, in a way similar to union.

If we want to use intersection on two queries of the same relation, we could just as easily set it up as a selection with a compound predicate.

```
(SELECT name, street, city, province, postal_code FROM  
owner_address)  
INTERSECT  
(SELECT name, street, city, province, postal_code FROM  
employee);
```

The **assignment** operation allows us to take the result of an expression and put it in a temporary variable.

The mathematical symbol for this is \leftarrow .

This is just notational convenience for the benefit of the reader.

So instead of something $\sigma_{owner_address_id=id}(license_plate \times owner_address)$, we could write:

$temp \leftarrow license_plate \times owner_address$

$\sigma_{owner_address_id=id}(temp)$.

SQL does have an assignment operator: `:=`.

I will, however, discourage you from using it.

The reason I discourage the use of assignment is the need to get away from the C-like thinking with counters and iteration.

The **Join** operations combines cartesian product with a selection.

The mathematical symbol for the *natural join* is \bowtie (bowtie).

It combines the cartesian product with a selection, forcing equality.

A natural join combined with a selection is a theta join \bowtie_{θ} where θ is the selection predicate.

Note that it is associative: so we can chain it as we need:

$vehicle \bowtie license_plate \bowtie owner_address$ is equivalent to
 $(vehicle \bowtie license_plate) \bowtie owner_address$ and
 $vehicle \bowtie (license_plate \bowtie owner_address)$.

Except, this probably doesn't work the way we hope on our sample data because the fields are not named the same thing.

If both fields are called `id`, the database server will try to match on those.

But this is not going to work for our data since they all have different names.

In SQL the keyword we need for this is `NATURAL JOIN`.

So we would do: `SELECT * FROM vehicle NATURAL JOIN license_plate;`

Again, though, this is not going to cooperate for us because our names don't match.

What do we do then?

The default type is the `INNER JOIN` and it's what you get if you just write `JOIN` instead of being specific.

It requires us to use the `ON` clause to specify how we relate one side to the other.

```
SELECT * FROM license_plate JOIN owner_address ON
owner_address_id = id; produces:
```

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2

Use the ON clause rather than NATURAL JOIN.

Oftentimes you need it and when you don't you sometimes get the wrong behaviour because a natural join does the wrong thing.

If the license plate relation called the plate number “id” instead of “number”, the natural join operation would run but produce no results.

Be explicit about how the join should work.

If there is an inner join, there is obviously an outer join. Actually, there are three!

Suppose we have a license plate that does not correspond to any address:

number	expiry	owner_address_id
ZZZZ 123	2018-09-30	24601
AAAA 855	2019-04-01	12949
BBCC 394	2019-02-12	null

and an address that does not correspond to any license plate:

id	name	street	city	province	postal_code
24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
86753	Sherlock Holmes	221B Baker St	London	ON	D4D 4D4

The inner join of license plate with owner address would look like this:

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2

If we use the left outer join, which has the symbol \bowtie ?

We get a tuple in the result for each entry in the left (first) table, even if it has no match in the second.

For such a tuple, all attributes from the right (second) table will simply be “null”.

If a tuple in the right table has no match in the left table, it does not appear at all.

The SQL for this is, of course, `LEFT OUTER JOIN`. So `SELECT * FROM license_plate LEFT OUTER JOIN owner_address ON owner_address_id = id;` returns:

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
BBCC 394	2019-02-12	null	null	null	null	null	null	null

The right outer join, symbol, \bowtie (RIGHT OUTER JOIN) is the mirror image operation of the left outer join.

We get a tuple in the result for each entry in the right (second) table, even if it has no match in the first.

For such a tuple, all attributes from the left (first) table will simply be “null”.

If a tuple in the left table has no match in the right table, it does not appear at all.

`SELECT * FROM license_plate RIGHT OUTER JOIN owner_address ON
owner_address_id = id;` returns:

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
null	null	null	86753	Sherlock Holmes	221B Baker St	London	ON	D4D 4D4

Finally, there is the full outer join, with the symbol \bowtie .

It does both the left and right outer join operations, including tuples from each relation that do not match the other:

number	expiry	owner_address_id	id	name	street	city	province	postal_code
ZZZZ 123	2018-09-30	24601	24601	Jean Valjean	19 Rue des Prisonniers	Ottawa	ON	B1B 1B1
AAAA 855	2019-04-01	12949	12949	Alice Jones	4 Generic Place	Kenora	ON	C2C 2C2
BBCC 394	2019-02-12	null	null	null	null	null	null	null
null	null	null	86753	Sherlock Holmes	221B Baker St	London	ON	D4D 4D4